

Postgres Performance for Humans

@craigkerstiens



Shameless plugs

<http://www.postgresweekly.com>

<http://www.craigkerstiens.com>

<http://www.postgresguide.com>

<http://www.postgresapp.com>

<http://www.heroku.com/postgres>

Postgres - TLDR

Postgres - TLDR

Datatypes

Conditional Indexes

Transactional DDL

Foreign Data Wrappers

Concurrent Index Creation

Extensions

Common Table Expressions

Fast Column Addition

Listen/Notify

Table Inheritance

Per Transaction sync replication

Window functions

NoSQL inside SQL

Momentum

TLDR in a quote

“It’s the emacs of databases”

<http://www.craigkerstiens.com/2012/04/30/why-postgres/>

OLTP vs OLAP

OLTP vs OLAP

Web apps

OLTP vs OLAP

BI/Reporting

Postgres Setup/Config

On Amazon

Use Heroku OR 'postgresql when its not your dayjob'


Other clouds

'postgresql when its not your dayjob'

Real hardware

High performance PostgreSQL

<http://thebuild.com/blog/2012/06/04/postgresql-when-its-not-your-job-at-djangocon-europe/>



Cache rules
everything around me



Cache Hit Rate

```
SELECT
    'index hit rate' as name,
    (sum(idx_blks_hit) - sum(idx_blks_read)) /
sum(idx_blks_hit + idx_blks_read) as ratio
FROM pg_statio_user_indexes
union all
SELECT
    'cache hit rate' as name,
    case sum(idx_blks_hit)
        when 0 then 'NaN'::numeric
        else to_char((sum(idx_blks_hit) -
sum(idx_blks_read)) / sum(idx_blks_hit + idx_blks_read),
'99.99')::numeric
    end as ratio
FROM pg_statio_user_indexes)
```

Cache Hit Rate

| name | ratio |
|----------------|-------|
| cache hit rate | 0.99 |

Index Hit Rate

```
SELECT
    relname,
    100 * idx_scan / (seq_scan + idx_scan),
    n_live_tup
FROM pg_stat_user_tables
ORDER BY n_live_tup DESC;
```

Index Hit Rate

| relname | percent_of_times_index_used | rows_in_table |
|---------------------|-----------------------------|---------------|
| events | 0 | 669917 |
| app_infos_user_info | 0 | 198218 |
| app_infos | 50 | 175640 |
| user_info | 3 | 46718 |
| rollouts | 0 | 34078 |
| favorites | 0 | 3059 |
| schema_migrations | 0 | 2 |
| authorizations | 0 | 0 |
| delayed_jobs | 23 | 0 |

Rough guidelines

Cache hit rate $\geq 99\%$

Index hit rate $\geq 95\%$

where on $> 10,000$ rows

Shortcuts

psql

```
$ cat ~/.psqlrc
```

```
\set ON_ERROR_ROLLBACK interactive
```

```
-- automatically switch between extended and normal  
\x auto
```

```
-- always show how long a query takes  
\timing
```

```
\set show_slow_queries
```

```
'SELECT  
    (total_time / 1000 / 60) as total_minutes,  
    (total_time/calls) as average_time, query  
FROM pg_stat_statements  
ORDER BY 1 DESC  
LIMIT 100;'
```

psql

```
$ cat ~/.psqlrc
```

```
\set ON_ERROR_ROLLBACK interactive
```

```
-- automatically switch between extended and normal  
\x auto
```



```
-- always show how long a query takes  
\timing
```

```
\set show_slow_queries
```

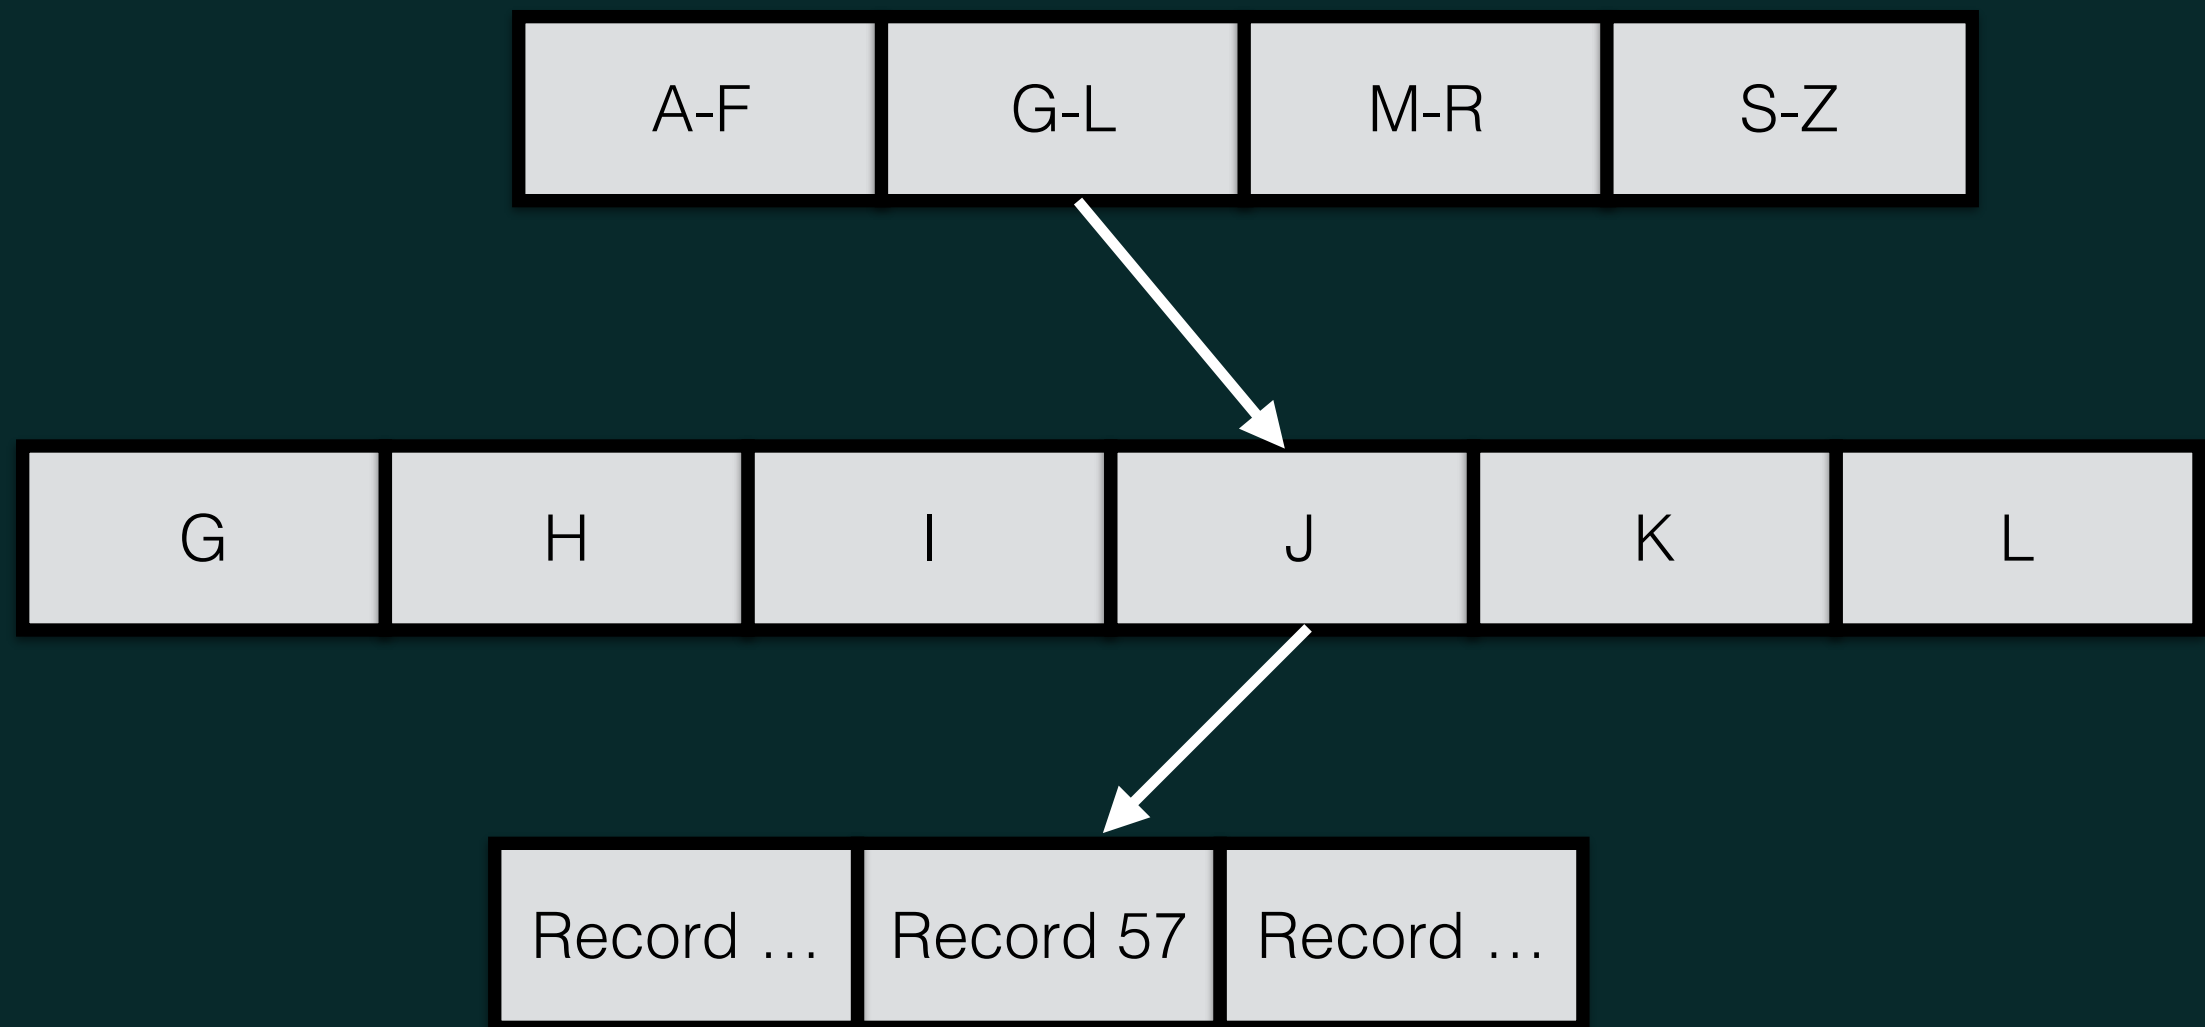
```
'SELECT  
    (total_time / 1000 / 60) as total_minutes,  
    (total_time/calls) as average_time, query  
FROM pg_stat_statements  
ORDER BY 1 DESC  
LIMIT 100;'
```

How Data is Retrieved

Sequential Scanning



Index Scans



Sequential Scans Index Scans

Good for large reports

Computing over lots of data (1k + rows)

Good for small results

Most common queries in your app

Understanding Specific Query Performance

Understanding Query Performance

```
SELECT last_name  
FROM employees  
WHERE salary >= 50000;
```

Explain

```
# EXPLAIN
  SELECT last_name
  FROM employees
 WHERE salary >= 50000;
```

QUERY PLAN

```
-----
Seq Scan on employees (cost=0.00..35811.00 rows=1
width=6)
  Filter: (salary >= 50000)
(3 rows)
```

Explain

```
# EXPLAIN
  SELECT last_name
  FROM employees
 WHERE salary >= 50000;
      QUERY PLAN
```

| | | | |
|--|--------------------------------------|--|-----------------------------------|
| Seq Scan on employees width=6) Filter: (salary >= 50000) (3 rows) | (cost=0.00..35811.00 startup time | | rows=1 max time rows return |
|--|--------------------------------------|--|-----------------------------------|

Explain Analyze

```
# EXPLAIN ANALYZE
SELECT last_name
FROM employees
WHERE salary >= 50000;
      QUERY PLAN
```

Seq Scan on employees (cost=0.00..35811.00 rows=1
width=6) (actual time=2.401..295.247 rows=1428
loops=1)

Filter: (salary >= 50000)

Total runtime: 295.379
(3 rows)

actual time

startup time

max time

rows return

Filter: (salary >= 50000)
(3 rows)

Rough guidelines

Page response times < 100 ms

Common queries < 10ms

Rare queries < 100ms

Explain Analyze

```
# EXPLAIN ANALYZE
SELECT last_name
FROM employees
WHERE salary >= 50000;
      QUERY PLAN
```

Seq Scan on employees (cost=0.00..35811.00 rows=1
width=6) (actual time=2.401..295.247 rows=1428
loops=1)

Filter: (salary >= 50000)

Total runtime: 295.379
(3 rows)

actual time

startup time

max time

rows return

Filter: (salary >= 50000)
(3 rows)

Indexes!

```
# CREATE INDEX idx_emps ON employees (salary);
```

Indexes!

```
EXPLAIN ANALYZE
  SELECT last_name
  FROM employees
 WHERE salary >= 50000;
          QUERY PLAN
```

```
-----
Index Scan using idx_emps on employees
(cost=0.00..8.49 rows=1 width=6) (actual time =
0.047..1.603 rows=1428 loops=1)
  Index Cond: (salary >= 50000)
Total runtime: 1.771 ms
(3 rows)
```


pg_stat_statements

```
$ select * from pg_stat_statements where query ~ 'from users where email';
```

| | |
|---------------------|--------------------------------------|
| userid | 16384 |
| dbid | 16388 |
| query | select * from users where email = ?; |
| calls | 2 |
| total_time | 0.000268 |
| rows | 2 |
| shared_blks_hit | 16 |
| shared_blks_read | 0 |
| shared_blks_dirtied | 0 |
| shared_blks_written | 0 |
| local_blks_hit | 0 |
| local_blks_read | 0 |
| local_blks_dirtied | 0 |
| local_blks_written | 0 |
| ... | |

pg_stat_statements

```
SELECT
    (total_time / 1000 / 60) as total,
    (total_time/calls) as avg,
    query
FROM pg_stat_statements
ORDER BY 1 DESC
LIMIT 100;
```

pg_stat_statements

| total | | avg | | query |
|--------|---|-------|---|-------------------------|
| ----- | + | ----- | + | ----- |
| 295.76 | | 10.13 | | SELECT id FROM users... |
| 219.13 | | 80.24 | | SELECT * FROM ... |

(2 rows)

Indexes

Indexes

B-Tree

Generalized Inverted Index (GIN)

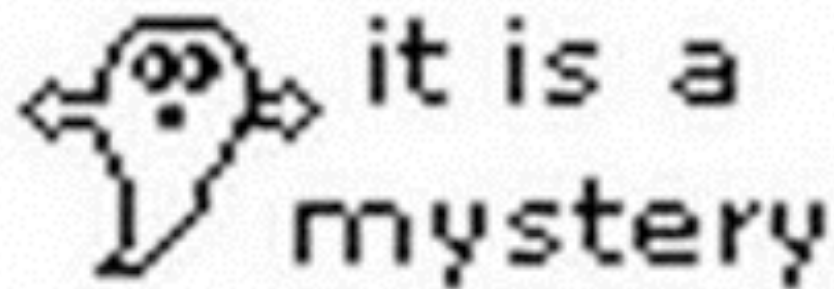
Generalized Search Tree (GIST)

K Nearest Neighbors (KNN)

Space Partitioned GIST (SP-GIST)

Indexes

Which do I use?



BTree

This is what you usually want

Generalized Inverted Index (GIN)

Use with multiple values in 1 column
Array/hStore

Generalized Search Tree (GIST)

Full text search

Shapes

Indexes

B-Tree

Generalized Inverted Index (GIN)

Generalized Search Tree (GIST)

K Nearest Neighbors (KNN)

Space Partitioned GIST (SP-GIST)

VODKA (Coming soon)

More indexes

Indexes

Conditional

Functional

Concurrent creation

Conditional

```
> SELECT *  
FROM places;
```

| name | population |
|-------|------------|
| ACMAR | 6055 |
| ARAB | 13650 |

Conditional

```
> SELECT *  
FROM places  
WHERE population > 10000;
```

| name | population |
|------|------------|
| ARAB | 13650 |

Conditional

```
> CREATE INDEX idx_large_population ON  
places(name) where population > 10000;
```

Functional

```
> SELECT *  
FROM places;
```

data

```
-----  
{"city": "ACMAR", "pop": 6055}  
{"city": "ARAB", "pop": 13650}
```


Functional

```
> SELECT *  
FROM places  
WHERE get_numeric('pop', data) > 10000;
```

data

```
-----  
{"city": "ARAB", "pop": 13650}
```

Functional

```
> CREATE INDEX idx_large_population ON  
places(get_numeric('pop', data));
```

Conditional and Functional

```
> CREATE INDEX idx_large_population ON  
places(data) WHERE  
get_numeric('pop', data) > 10000;
```

One more thing

CREATE INDEX **CONCURRENTLY** ...

roughly 2-3x slower

Doesn't lock table

history / JSON

hstore

```
CREATE EXTENSION hstore;  
CREATE TABLE users (  
    id integer NOT NULL,  
    email character varying(255),  
    data hstore,  
    created_at timestamp without time zone,  
    last_login timestamp without time zone  
);
```

hstore

```
INSERT INTO users
VALUES (
  1,
  'craig.kerstiens@gmail.com',
  'sex => "M", state => "California"',
  now(),
  now()
);
```

JSON

SELECT

```
'{"id":1,"email":  
  "craig.kerstiens@gmail.com",}'::json;
```


hstore

Indexes work

gin

gist

json

Functional indexes work
have fun

jsonb

The world is better

Pooling

Options

Application/Framework layer

Stand alone daemon

PG options

pgbouncer

pgpool

Adding Cache

Replication options

slony

londiste

bucardo

pgpool

wal-e

barman

Replication options

slony

londiste

bucardo

pgpool

wal-e

barman

Backups

Logical

pg_dump

can be human readable, is portable

Physical

The bytes on disk

Base backup

Logical

Good across architectures
Good for portability

Has load on DB

Works < 50 GB

Physical

More initial setup
Less portability

Limited load on system

Use above 50 GB

Recap

Recap

OLAP

Whole other talk

Disk IO is important

Order on disk is helpful (pg-reorg)

MPP solutions on top of Postgres

Recap

OLTP (webapps)

Ensure bulk of data is cache

Optimize overall query load with `pg_stat_statements`

Efficient use of indexes

When cache sucks, throw more at it

Questions

[http://www.speakerdeck.com/u/
craigkerstiens/](http://www.speakerdeck.com/u/craigkerstiens/)